

Home Equity in Retirement*

Makoto Nakajima[†]
Federal Reserve Bank of Philadelphia

Irina A. Telyukova[‡]
Mulligan Funding

October 5, 2018

Abstract

Retired homeowners dissave more slowly than renters, which suggests that homeownership affects retirees' saving decisions. We investigate empirically and theoretically the life-cycle patterns of homeownership, housing and nonhousing assets in retirement. Using an estimated structural model of saving and housing decisions, we find, first, that homeowners dissave slowly because they prefer to stay in their house as long as possible, but cannot easily borrow against it. Second, the 1996-2006 housing boom significantly increased homeowners' assets. These channels are quantitatively significant; without considering homeownership, retirees' net worth would be 28-53% lower, depending on age.

JEL classification: D91, J26, E21, G11

Keywords: Housing, Retirement Saving Puzzle, Mortgage, Health, Life-cycle

*We thank Eric French for his generous help with the HRS Exit Waves, and Orazio Attanasio for several discussions about this project. For their comments and suggestions, we also thank Jonathan Heathcote and Gianluca Violante, and the participants of seminars at NYU, Fordham University, Cambridge University, Bristol University, University of Warwick, University of Mannheim, Uppsala University, Wisconsin Business School, University of Hawaii, Stony Brook University, University College London, Columbia University, Federal Reserve Banks of Chicago, Atlanta, San Francisco, New York and St. Louis, CEMFI, UCSD, USC Lusk School, and of the 2012 NBER-Oxford Said-EIEF-CFS Household Finance Conference, 2012 Regensburg Real Estate Conference, 2011 REDg VI meetings, 2010 Econometric Society World Congress, CEF Meetings, NBER Summer Institute EFACR group, UCSB-LAEF Credit, Default, and Bankruptcy Conference, and 2009 UW-Atlanta Fed Housing-Labor-Macro-Urban Conference. We acknowledge financial support of Networks Financial Institute at Indiana State University via its award in the 2011 Financial Services Regulatory Reform Paper Competition, as well as of the UCSD Faculty Career Development Grant. The views expressed here are those of the authors and do not necessarily represent the views of the Federal Reserve Bank of Philadelphia or the Federal Reserve System. Previous versions of the paper were circulated with the title "Home Equity Withdrawal in Retirement". Last but not least, we thank Walter for inspiration.

[†]Research Department, Federal Reserve Bank of Philadelphia. Ten Independence Mall, Philadelphia, PA 19106-1574. E-mail: makoto.nakajima@phil.frb.org.

[‡]Mulligan Funding. 4619 Viewridge Avenue, San Diego, CA 92123-5611. E-mail: irina.a.telyukova@gmail.com.




Figure 1: Median Life-Cycle Wealth Profiles. Source: HRS.




Figure 2: Normalized Median Life-Cycle Wealth Profiles. Source: HRS

1 Introduction

An important question in the life-cycle savings literature is why the elderly dissave slowly in the data. As figure 1 shows over the period 1996-2006, median net wealth remains high very late into the life cycle. The observation that many people die with significant savings, which is puzzling in the context of a simple life-cycle model, has been termed the “retirement saving puzzle”.

The picture changes dramatically if we consider the saving behavior of retirees who own homes, compared to those who do not. Consider figure 2, which documents cohort profiles of median net worth over the same period, normalized by the first observation, for homeowners versus renters.¹ The difference is stark. Homeowners have flat or increasing profiles of net wealth over this period, while renters display a far faster rate of asset decumulation. This suggests that housing may play a major role in determining how retirees save or dissave.

Motivated by this observation, in this paper we examine the role of housing in retirees’ saving behavior. Rather than explaining only the life-cycle profile of household net worth, as the literature has done, we seek to understand several facts about saving in retirement concurrently: we consider both housing and nonhousing assets, as well as homeownership rates and collateralized debt. Our main question is what role housing plays in accounting for the retirement saving puzzle. We find that considering explicitly the nature of housing as a complex asset with properties different from other assets makes an important difference in understanding retiree saving behavior, and changes our conclusions regarding the retirement saving puzzle, relative to previous literature. By understanding the nature of homeownership late in life, we also shed new

¹In this figure, homeowners are the households that start in our 1996-2006 HRS sample as owners, and remain so throughout the sample; renters are similarly “perpetual” renters.

light on the role of bequest motives, uncertainty and precautionary motives in retirement.

We begin by documenting in detail, using the Health and Retirement Study (HRS) over the period 1996-2006, various facts about retirees' financial and housing asset holdings, and about their use of home equity. Because the HRS is a longitudinal survey, it allows us to study life-cycle asset and debt profiles over time. We also document other relevant changes in retirees' lives pertaining to their income, health, medical expenses, marital status and the like, which are all potential drivers of saving decisions late in life.

We then build a model of household saving decisions in retirement with the goal of accounting for the life-cycle facts of interest. The model includes both financial assets and a house, which serves as an asset but also provides utility. Retirees can choose whether to own a home or rent, and homeowners can access their home equity by selling the house or by secured borrowing, with an age-dependent borrowing constraint. Retirees have a warm-glow bequest motive, and face idiosyncratic uncertainty in their health status, medical expenses, and longevity, as well as that of their spouse. They have Social Security and pension income, and access to a government-provided social insurance program which provides a consumption floor for households who run out of assets. Aggregate house price dynamics reflect recent changes in the U.S. housing market.

The key potential driving forces of retiree saving behavior in the model are nonfinancial and financial benefits of homeownership, including the housing price boom of 1996-2006, differential liquidity properties of housing and financial assets, as well as bequest motives, longevity risk, medical expense risk and the Medicaid-like government social safety net. While bequest motives, longevity, medical expense risk and Medicaid have been studied and debated in previous literature on the retirement saving puzzle, the housing-related forces have not been considered previously in the context of a structural model; yet, they may matter on their own, and may interact with bequest motives and health risk. We estimate our model, and use it to quantify the role of each of these forces, as well as to understand their interactions.

We estimate the model by a two-step procedure using the HRS. We measure exogenous parameters, such as the shock processes, outside the model, then use the model to estimate other parameters by a minimum-distance estimator, targeting jointly the life-cycle facts mentioned before. Our model successfully replicates these facts, with reasonable resulting parameter values.

To understand the quantitative contribution of the key model features, we conduct a series of experiments using the benchmark model. First, we shut down these mechanisms one at a time, keeping the rest of the model unchanged. Then, we strip the model down to the basic life-cycle framework, and introduce the main features in sequence in different orders, to understand and quantify the interactions between them.

We find that the high homeownership rate late into the life-cycle that we observe in the data is crucial to consider for understanding retiree saving behavior. Housing-related channels are significant contributors to the retirement saving puzzle. Retirees stay homeowners late in life, due to financial benefits and attachment to their homes, but become increasingly locked into their home equity; we find that borrowing constraints on retirees tighten considerably as they age. This means, on the one hand, that those who remain homeowners do not decumulate their home equity, thus creating the kind of flat housing and net worth profile that we see in the data, while those who face a large expense may come up against their borrowing constraint and be forced to sell the house. In addition, those who owned a house in the period 1996-2006 became beneficiaries of the housing boom, which further contributes to the flat or increasing net worth profiles of elderly homeowners. These effects together are a big part of what creates the stark difference that we observe between homeowners and renters in the data. The effects are robust to perturbation of the parameters of the model, as we show in sensitivity experiments.

We also use the model to understand the relative importance of bequest motives and medical expense risk. In our benchmark estimated model, we show that bequest motives play a quantitatively significant role in accounting for the retirement saving puzzle, contributing between 8 and 21% of median net worth depending on age, while the role of medical expenses and expense risk is moderate, at 2-14%. We also show in sensitivity experiments, however, that the strength of the bequest motive and the precautionary motive are difficult to identify separately in this class of models, so that it is possible for medical expenses to play a larger quantitative role, possibly, though not necessarily, at the expense of the contribution of bequest motives. This points to the need for a richer model of bequest motives in future research. It is still notable that the role of housing is never diminished by changes in relative roles of risk and bequests.

In addition to the quantitative decompositions, we conduct an experiment where we allow households to make a decision on whether or not to maintain their home. We evaluate this as an additional, possibly hidden, channel of asset decumulation, consistent with data evidence that homes of elderly owners depreciate more quickly than those of younger owners (Davidoff (2006)). We treat this as a hidden channel because we assume that self-reported housing values of owners who remain in their houses do not take into account the depreciation rate unless they have the house appraised for sale, for example. We find this to be a significant channel of asset decumulation that increases in importance with age. On average, about 15% of model homeowners choose not to maintain their homes in the 75-85 year old cohort; by age 95, nearly 40% of homeowners choose not to maintain the house.

We thus have three main contributions. First, our careful documentation of the longitudinal

data provides a set of facts regarding retirees’ saving behavior in more detail than previously studied. In addition to being of empirical interest, we think it is important that these facts should be considered explicitly by a theory that seeks to explain saving behavior in retirement. Second, our model enables us to describe the tradeoff between housing and nonhousing assets in retirement, and to characterize the reasons for homeownership in retirement. To our knowledge, we are the first to do this in the context of a rich structural model. Third, we address the retirement saving puzzle from a new perspective.

The remainder of the paper is organized as follows. Section 2 discusses related literature. Section 3 describes our data and stylized facts. Section 4 develops the model. Section 5 describes the estimation strategy, presents the resulting parameters and assesses the fit of the model. Section 6 describes the quantitative decompositions that we perform in our model. Section 7 discusses robustness of key results and identification of parameters, based on sensitivity analysis. Section 8 describes the experiment of endogenizing the home maintenance decision. Section 9 concludes. Some details of data analysis and quantitative experiments are in the appendix.

2 Related Literature

Our paper is related to a number of studies of savings decisions in retirement. On the retirement saving puzzle itself, several answers have been proposed. Hurd (1989) estimates the life-cycle model with mortality risk and bequest motives, and finds that intended bequests are small. Love et al. (2009) analyze the retirement saving puzzle using “annualized comprehensive wealth,” which is a measure of total wealth, including annuity-like assets as well as financial and nonfinancial assets. Hubbard et al. (1995) argue that means-tested social insurance programs provide a virtual consumption floor and thus strong incentives for low-income individuals not to save; their paper can thus be seen as reinforcing the retirement saving puzzle. Ameriks et al. (2011) study the relative importance of bequest motives and public care aversion in explaining the related annuity puzzle using a model of retirement and survey data, and find both motives significant in the data. Lockwood (2012) looks at the related decision of whether or not to purchase life insurance in retirement.

The recent paper by De Nardi et al. (2010) is most closely related to ours in terms of approach. They estimate a life-cycle model of retirees using the AHEAD sub-sample of the HRS, focusing on singles among the oldest old. Like them, we use a life-cycle model of retirees together with the HRS, with health and medical expenditures being a major source of uncertainty for retirees. The key difference between our work and theirs is our focus on housing and home equity borrowing, while they aggregate all the assets in the household portfolio, and study the

profile of the consolidated asset position in retirement. In addition, we consider in our sample both single and couple households, while De Nardi et al. (2010) include only singles.

The empirical part of our paper is related to Venti and Wise (2004), one of whose main findings, confirmed by our data analysis, is that retirees rarely downsize their houses even in older age, unless a drastic event such as illness or death of a spouse occurs. They also provide evidence from the HRS that some older households move into larger homes; we show that this mostly appears to be the case based on rising house prices, rather than reflecting purchases of larger homes, a possibility pointed out by Skinner (2004).

Other studies of implications of health and medical expenditure risks on portfolio decisions of retirees is Yogo (2016), which treats health expenses as endogenous investment, and Kopecky and Koreshkova (2014), who focus on nursing home expenses and study the implications on aggregate savings and the wealth distribution. Marshall et al. (2010) revisit the measurement of end-of-life medical expenses in an empirical exercise, and find these expenses to be significant.

More generally, our paper fits with the recently growing body of work that incorporates housing explicitly into a macroeconomic framework. For example, Gervais (2002) studies the effects of the preferential tax treatment of housing on capital accumulation. Fernández-Villaverde and Krueger (2011) use a general equilibrium life-cycle model to study the life-cycle profile of housing and nonhousing consumption. Díaz and Luengo-Prado (2010) investigate housing as an explanation for observed large wealth inequality in the U.S.²

3 Facts

We begin by describing the most relevant data facts about homeownership and saving in retirement, which we want to explain using our theory. In addition to the facts already presented, these are retirees' life-cycle profiles of homeownership rates, housing and nonhousing assets, the rate of collateralized debt and the amount of debt held. We also present some facts that inform our modeling choices, as we describe below. We then give much more detail on the mapping between the data and the model in section 5.

3.1 Data

The Health and Retirement Study (HRS) is a biennial longitudinal survey of households of age 50 and above, conducted by the University of Michigan. The survey began in 1992. Due to

²Other studies of housing that use structural models include Davis and Heathcote (2005), who study housing in a business cycle model, and Chambers et al. (2009), who investigate the optimal choice between different types of mortgages in a general equilibrium model. Ortalo-Magné and Rady (2006) study the impact of income shocks and credit constraints for business cycle dynamics of the housing market.

issues with the early data on assets (see De Nardi et al. (2010)), we begin our data observation in 1996 and use six waves that span 10 years, through 2006. We use the RAND version of the HRS data, constructing a full merged set from the flat files provided by RAND; in addition, we merge in information from the exit waves of the survey, concerning members of the sample who die between two waves, in order to accurately measure medical expenses at the end of life.

We consider everyone present in the 1996 sample who is of age 63 and above and who reports being retired, either fully or partially. We consider both couples and single households. We subdivide the sample into six cohorts, of ages 63-67, 68-72, 73-77, 78-82, 83-87, and 88-97 in 1996. We follow these cohorts across the waves of the survey and document their life-cycle patterns of asset holding and health, as described below. Because assets are measured in the HRS at household level, while health and other demographics are at the individual level, we adjust the weighting schemes appropriately to construct information for our model households.

The HRS sample is replenished several times over the course of the survey. There are multiple ways to deal with this cohort replenishment: one could only consider those who appear in the survey starting in 1996, or include in later waves everyone who belongs to a given cohort by age, even if they enter the survey after 1996. As a benchmark, we consider only households that appear in the 1996 wave, without cohort replenishment. For robustness, we have considered an alternative in which we replenish the cohorts after the 1996 wave; see appendix Appendix A.

A related issue with the HRS sample is weighting. Each individual in the HRS is assigned a wave-specific weight each year he appears in the sample; however, an individual who lives in a nursing home is assigned a weight of zero. We do not want to lose such individuals from the sample. In order to compute weighted statistics, but not lose nursing home residents, we apply the weight attached to each individual in the initial (1996) wave of our sample, consistent with our choice of unreplenished cohorts. For robustness, we reconstructed all of our analysis with the replenished sample, using the weights specific to each wave; we also constructed unweighted measures, for the purpose of comparing with De Nardi et al. (2010). We discuss these measures in appendix Appendix A. Our choices imply that we consider an unbalanced panel in our analysis, since households will drop from the sample due to mortality. This choice is the most consistent with our model, where we will construct an equivalent unbalanced panel, with households dying according to the same probability as in the data. We will discuss this more in the Estimation section.

To allow our data measures to map into the model, we measure financial assets as the sum of non-housing assets (excluding businesses and cars) net of all debt, including home equity debt. We track housing assets separately, including only the primary residence, since other real estate




Figure 3: Normalized Median Net Worth, Perpetual Homeowners versus Switchers

information is not available in all waves of the survey. Finally, we define total assets as the sum of financial and housing assets, net of all debt.³ Our definition of nonfinancial income includes Social Security, pension, disability, annuity, and government transfer income. Because some of our retirees are only partly retired, we also include labor income in this measure; overall, labor income plays a small role, constituting on average about 6% of total income.

As homeowners in our data, we take everyone who reports owning their residence. In the other category, labeled “renters”, we include not only actual renters, but also individuals living in nursing homes, with their children, and in other arrangements not involving homeownership. The results presented here are robust to this aggregation of non-owners. A few of the nursing home residents report owning a home. For such individuals, we set the value of their home to zero, and fold their house value into their financial assets.

3.2 Life-cycle Profiles

To accompany our motivating figure 2, we want to confirm that staying a perpetual homeowner throughout the sample implies different saving behavior than that of anyone who becomes a renter any time in the sample, and that the difference that we observe is not simply a function of being in different wealth quintiles. In figure 3, we plot the normalized median net worth profiles of those who are homeowners perpetually versus those who switch into renting at some point in the sample. We designate them by the wave in which they switch, and present only three cohorts to make the figure legible. This figure confirms that when one sells the house, one decumulates assets more quickly than if one stays in the house, and that this behavior is not simply a function of overall wealth; thus, we need to consider housing separately from other assets.

³We experimented with other definitions of assets and found that the results are not affected.




Figure 4: Homeownership Rates.




Figure 6: Median Housing Assets, Conditional on Ownership.




Figure 5: Median Total Assets.




Figure 7: Median Financial Assets.

Figure 4 shows the life-cycle profile of homeownership rates among retirees (dark blue solid line). In general, homeownership rates are declining with age, from around 90% at age 65 to just below 40% by age 95. We also break down the rates by the size of the household. The breakdown shows that conditional on household size, the decline is milder than the overall average for 2-adult households, suggesting that the overall decline in homeownership may be driven in part by a transition from a 2-person to a 1-person household. This agrees with the findings of Venti and Wise (2004). Motivated by this finding, we will allow our model households to change size over time, according to probabilities consistent with the data.

Figure 5 again plots the life-cycle median net worth profiles of retirees. Figures 6 and 7 break down these profiles into housing and financial assets, as defined above. Total asset holdings are increasing with age for the youngest three cohorts, while they are slightly declining for the older cohorts. The breakdown into housing and non-housing assets shows that the increase in total asset value for the younger cohorts is mainly driven by increasing housing assets, while financial




Figure 8: Proportion of Households in Net Debt.




Figure 9: Median Net Debt Holding Among Net Debtors.




Figure 10: Proportion of Households with Gross Secured and Unsecured Debt.




Figure 11: Median Gross Secured and Unsecured Debt among Debt Holders.

assets are relatively flat for each cohort. Note that the median financial asset profile includes households who sell their house after 1996 and convert it into financial assets; this composition effect is likely important in keeping financial asset profiles flat rather than decreasing.

Figures 8 and 9 plot the shares of retirees who are in debt by our model definition, that is, those who hold a negative financial asset position, as well as the median amount of debt held, conditional on being a debtor. The share of debtors is decreasing with age, from around 18% at age 65 for the first cohort, to nearly zero for the oldest cohort. The conditional debt profile is increasing for the younger cohorts, reflecting that those with more debt stay in debt longer.

To understand how debt should be modeled, we also consider the profiles of *gross* secured and unsecured debt. The proportion of households with each type of debt in figure 10 decreases with age, in a fashion similar to the negative financial asset position; slightly fewer retired

households have unsecured debt than secured debt. The of conditional debt levels in figure 11 replicate the median net debt profile above. Instead, the amount of unsecured debt (right scale) is relatively small, at maximum \$2,000 for the youngest cohort, compared to \$30,000-40,000 in secured debt, decreasing over the lifecycle, and approximately flat for each cohort. Due to low median unsecured debt, and to reduce the computational burden, we will assume unsecured debt away in the model, so that those without a house will not be able to borrow.

4 Model

We focus on retiree households, which allows us to abstract from labor supply and retirement decisions. A household in the model starts out either single or as a couple; couple households can become single if one spouse dies, but single households do not re-marry. This assumption is motivated by the data, where the number of remarriages in retirement is small.

A household starts out as a homeowner or a renter. In each period, the household chooses consumption and financial saving, and makes a decision regarding housing. For a homeowner, the housing decision is whether to move out of the house or to stay in it. Homeownership provides utility benefits, in addition to consumption services from the house; these capture factors such as attachment to one's house and neighborhood, the ability to modify one's house to individual taste, but also some financial benefits of ownership that are not explicitly in the model, such as tax exemption of imputed rents of owner-occupied housing, mortgage interest deduction, or insurance against rental rate fluctuation. In addition, homeowners are able to borrow against their home equity; the collateral constraint can change with age, as discussed below. For a renter, the housing choice is only the size of the rental property. We abstract from the decision of a homeowner to move to a different house, or the decision of a renter to buy a house. These abstractions are made to simplify the problem, but are motivated by the observation in the data that the proportion of homeowners making downsizing moves is small, as is the proportion of renters who purchase a home late in life. Finally, renters are not able to borrow due to low unsecured debt among retirees in the data.

The aggregate price of housing in the model is increasing to capture the housing boom of 1996-2006. We assume that households anticipate the increase in a deterministic fashion, and do not face idiosyncratic house price shocks. This last assumption is necessary, given the complexity of the problem. In addition to the household size shock, households are subject to two other types of idiosyncratic shocks: health status, conditioned on age, which includes the probability of death, and out-of-pocket medical expenditures, conditioned on age, size, income and health.

In addition to financial asset income, households have access to pension income. Since in

the data nonfinancial income is stable over time conditional on household size, in the model we assume income time-invariant, as long as household size does not change. In addition, households have access to a government-provided consumption floor, which captures insurance programs for the elderly such as Medicaid. Finally, households have a warm-glow bequest motive.

4.1 Preferences

A household is born as a retiree at model age $i = 1$. The household potentially lives up to age I , but dies stochastically; this is discussed more below. The household maximizes its life-time utility. The utility function is time-separable with subjective discount factor β . The period utility function has the following form:

$$u(c, h, s, o) = s \frac{\left(\frac{1}{\mu_s} c^\eta (\omega_o h)^{1-\eta} \right)^{1-\sigma}}{1-\sigma} \quad (1)$$

where c is nonhousing consumption, h is consumption of housing services, $s \in \{1, 2\}$ is the number of adults in the household, and $o \in \{0, 1\}$ is the tenure status, with $o = 0$ representing renting, and $o = 1$ representing owning. We assume a linear technology from the size of the house to the quantity of housing services, which implies that h is the size of the house that the household lives in as well. Consumption is aggregated by a Cobb-Douglas function, with η determining the relative importance of housing and nonhousing consumption. The period utility function applied to the aggregated goods is a standard CRRA function with risk aversion parameter σ . μ_s is the household equivalence scale conditional on household size.⁴ In particular, if $\mu_1 = 1$ and $\mu_2 \in (1, 2)$, the *household-size multiplier* for a one-adult household is $\frac{1}{\mu_1^{1-\sigma}} = 1$, while the multiplier for a two-adult household is $\frac{2}{\mu_2^{1-\sigma}} > 1$ for $\sigma > 0$. In other words, the assumption captures the benefits of having multiple adults instead of one adult in the household. ω_o captures the extra utility from owning a house. We normalize the renters' $\omega_0 = 1$.

As in De Nardi et al. (2010), a household gains utility from leaving bequests.⁵ When a household dies with consolidated wealth of a , the household's utility function takes the form:

$$v(a) = \gamma \frac{(a + \zeta)^{1-\sigma}}{1-\sigma}. \quad (2)$$

⁴For a more detailed discussion on the household equivalence scale, see Fernández-Villaverde and Krueger (2007). Li and Yao (2007) make a similar assumption with respect to the effect on the household size on utility.

⁵De Nardi (2004) finds that the bequest motive is important in capturing the observed wealth distribution, especially the wealth concentration, using a general equilibrium overlapping-generations model with accidental and intended bequests.

Here, γ captures the strength of the bequest motive, and ζ affects marginal utility of bequests.

4.2 Household Structure, Health and Mortality

Households in the model are distinguished demographically in terms of their size and health. The health status of a household is represented by $m \in \{0, 1, 2, \dots, M\}$, where $m = 0$ represents death, which is an absorbing state, i.e. $m_j = 0$ for $\forall j \geq i$ if $m_i = 0$. $m > 0$ indicates that the household is alive and in one of several time-varying health states. We assume that m follows a first-order Markov process. $\pi_{i,m,m'}^m$ is the transition probability from health m to m' for an age- i household, and it includes survival probability of households.

$s \in \{1, 2\}$ represents the number of adults in a household. The transition from $s = 2$ to $s = 1$ can capture the death of a spouse or a divorce; in our estimation, we will abstract from divorces and remarriages, as we find these to be rare in the data. Thus, one-adult households ($s = 1$) remain single for the rest of their life. In contrast, two-adult households ($s = 2$) stochastically change to one-adult households. Household size transition probabilities are age-dependent and denoted by $\pi_{i,s,s'}^s$. By assumption, $\pi_{i,1,1}^s = 1$, $\pi_{i,1,2}^s = 0$ for all i .

Household size thus affects household decisions as follows. First, two-adult households maximize the sum of the utilities of the two. In order to avoid keeping track of types of each member of a couple, we assume that the two adults have the same utility function, so the utility of a couple is that of a one-adult household multiplied by two, as captured by s in the utility function. Second, consumption is split equally in two-adult households. However, each of the household members enjoys more than half of the consumption because of the positive externality within the household. This is captured by μ_s in the utility function. Third, pension income depends on household size. Finally, two-adult households face the size shock. Together with the mortality shock embedded in $\pi_{i,m,0}^m$, this means that in a couple, one spouse can die first via the stochastic shock to s , or both spouses can die at the same time via the household-wide mortality shock.

4.3 Medical Expenditures

Household health introduced above has two effects. First, survival probability is lower for a household in worse health; second, out-of-pocket medical expenses are on average higher for a household in worse health. Both are facts from our data (see section 5). A household is hit by out-of-pocket (uninsurable) medical expenditure shocks $x \in \{x_0 = 0, x_1, x_2, \dots, x_X\}$, which are a function of its age i , health m , income b and size s . The probability that a given x is drawn is denoted by $\pi_{i,m,b,s}^x$. Conditional on age, size, income and health, medical expense shocks are i.i.d.; however, because these household characteristics are persistent, so are medical expenses.

4.4 Nonfinancial Income

We assume that the household's nonfinancial income is $\psi_s b$, where $b \in \{b_1, b_2, b_3, \dots, b_B\}$ and ψ_s adjusts the nonfinancial income according to the number of adults in the household, with $\psi_1 = 1$. b is different across households, but is time-invariant for each household. This assumption captures the fact that the main sources of income for retirees are Social Security and other pension benefits, which are fixed at time of retirement and appear constant thereafter in our data.

4.5 Housing

A household is either a renter ($o = 0$) or a homeowner ($o = 1$). A homeowner with a house $h \in \{h_1, h_2, h_3, \dots, h_H\}$ decides whether to move out of the house and become a renter, or to stay in the same house. The total value of the house is $p_1 h$, where p_1 is the current house price for owners. If a homeowner sells her house, she receives its value net of any debt, from which she pays a proportional cost κ of moving out, and a capital gains tax. In addition, the homeowner has to pay a proportional maintenance cost δ while she lives in the house. In the benchmark version of the model, we assume that everyone pays this cost, which we will relax later.

The house price p_1 is assumed to have only an aggregate time-varying component; we do not consider heterogeneity of house prices, in order to keep the problem manageable. We further assume that households expect house prices to grow at a constant rate g_1 , consistent with the upward price trend in 1996-2006. As an alternative, we have tried the assumption that households expect house prices to stay constant, treating all growth in house prices from the exogenous price trend as a surprise. These two alternatives yield nearly identical results in terms of household behavior; we choose the former specification as it is consistent with rational expectations.

A renter chooses the size of the rental property h each period. Unlike owners, renters can move between properties of different sizes at no moving cost. All rental contracts are for one period. The per-period rental rate is r_h , which consists of two components: $r_h = r + \delta$. Here, r is the riskless interest rate, discussed more below. The rental rate captures the competitive cost to an intermediating real estate firm of holding housing and renting it out.⁶

Rental properties are evaluated at price p_0 , which grows deterministically at a constant rate g_0 . This allows us to capture the effect of the 1996-2006 housing boom on the rental market. As dictated by the data, $g_0 < g_1$.

⁶See Nakajima (2010) for a more detailed discussion about the determination of the rental rate.

4.6 Saving and Home Equity Borrowing

We use a to denote the household's consolidated financial asset balance. Households can save ($a > 0$) at interest rate r . In addition, home equity borrowing is allowed; homeowners can borrow against the value of their house at the rate of $r + \xi$, where ξ is the mortgage premium. The borrowing limit in period t has the form $a \geq -(1 - \lambda_i)hp_1$: a homeowner can borrow up to a fraction $1 - \lambda_i$ of the value of the house (hp_1) in each period. While the parameter λ_i can most directly be interpreted as a downpayment constraint, in this setup we are agnostic about the exact type of equity loan contracts available and the associated cost types. Therefore, we intend for it to capture parsimoniously all direct costs of borrowing against home equity, e.g. the costs of refinancing, the costs of opening a new home equity line of credit (HELOC), or the upfront costs of a reverse mortgage. We allow this parameter to be age-specific. While there are no overt age requirements for traditional mortgage loans that we are aware of, Caplin (2002) points out that many retired homeowners cannot qualify for conventional mortgages because they fail income requirements of such loans. Our specification can capture such age variation in borrowing constraints. We will estimate the parameters λ_i from the model. As we mentioned above, renters in the model cannot borrow, because there is little unsecured debt among retirees in the data.

4.7 Government Transfers

Following Hubbard et al. (1995) and De Nardi et al. (2010), we assume that the government uses means-tested social insurance, which effectively provides a consumption floor. This is especially important in our model because a large out-of-pocket medical expenditure shock could force a household to negative consumption in the absence of social insurance. The per-adult consumption floor is denoted by \underline{c} . Following De Nardi et al. (2010), we assume that the government subsidizes each member of a household up to \underline{c} if the household runs down its financial assets. Homeowners are eligible for the consumption floor so long as the value of their house is below some threshold $\bar{h}p_1$, which captures the Medicaid homestead exemption (De Nardi et al. (2016)).

4.8 Household Problem

We formalize the household problem recursively, and separately for homeowners and renters. We use a prime to denote a variable in the next period. The state variables of a household are (i, s, b, m, x, p, h, a) : age, size, income, health status, medical expenses, house price, amount of housing, and its financial assets. We use $h = 0$ to represent a renter; $h > 0$ means that a household is a homeowner with a house size of h . House prices are a vector $p = (p_0, p_1)$.

Beginning with the problem of the renter, the Bellman equation is:

$$V(i, s, b, m, x, p, 0, a) = \max_{\tilde{h}, a' \geq 0} \left\{ u(c, \tilde{h}, s, 0) + \beta \sum_{s'} \pi_{s,s'}^s \sum_{m' > 0} \pi_{i,m,m'}^m \sum_{x'} \pi_{i+1,m',b,s',x'}^x V(i+1, s', b, m', x', p', 0, a') + \beta \pi_{m,0}^m v(a') \right\} \quad (3)$$

subject to:

$$\tilde{c} + a' + r_h \tilde{h} p_0 + x = (1+r)a + \psi_s b \quad (4)$$

$$c = \begin{cases} \max\{s\underline{c} - r_h \tilde{h} p_0, \tilde{c}\} & \text{if } a' = 0 \\ \tilde{c} & \text{otherwise} \end{cases} \quad (5)$$

$$p'_i = (1 + g_i) p_i \text{ for } i = \{0, 1\} \quad (6)$$

The renter chooses the level of assets to carry over to the next period (a') and the property that he rents in the current period (\tilde{h}) to maximize the sum of three components. The first is the period utility. The second is the discounted expected future value conditional on surviving in the next period ($m' > 0$), with the expectation based on the transition probabilities for the household size, health and medical expense shocks. Income b is constant, and the renter remains a renter ($h' = h = 0$). The third component of (3) is the utility from bequests, in case of death ($m' = 0$). Notice that, for a renter, the only assets left as estate are the financial assets (a'). Equation (4) is the budget constraint, while (5) is the government-provided consumption floor, net of rental expenses, available only when the renter runs down financial assets ($a' = 0$).

A homeowner chooses between staying in his house (V_1) and selling to become a renter (V_0):

$$V(i, s, b, m, x, p, h, a) = \max\{V_0(i, s, b, m, x, p, h, a), V_1(i, s, b, m, x, p, h, a)\} \quad (7)$$

A homeowner who decides to sell the house and become a renter solves:

$$V_0(i, s, b, m, x, p, h, a) = \max_{a' \geq 0} \left\{ u(c, h, s, 1) + \beta \sum_{s'} \pi_{s,s'}^s \sum_{m' > 0} \pi_{i,m,m'}^m \sum_{x'} \pi_{i+1,m',b,s',x'}^x V(i+1, s', b, m', x', p', 0, a') + \beta \pi_{m,0}^m v(a') \right\} \quad (8)$$

subject to (6) and:

$$\tilde{c} + a' + x + (\kappa + \delta)hp_1 + q(s, h, p_1) = hp_1 + (1 + \tilde{r})a + \psi_s b \quad (9)$$

$$q(s, h, p_1) = \tau_q(hp_1(1 - \kappa) - h\bar{p}_1 - \bar{q}s) \quad (10)$$

$$c = \begin{cases} \max\{s\underline{c}, \tilde{c}\} & \text{if } a' = 0 \\ \tilde{c} & \text{otherwise} \end{cases} \quad (11)$$

$$\tilde{r} = \begin{cases} r & \text{if } a' \geq 0 \\ r + \xi & \text{if } a' < 0 \end{cases} \quad (12)$$

There are four differences from the renter's problem shown above. First, the current tenure status is a homeowner ($o = 1$) with the house size of h , as can be seen in the period utility function. Second, the budget constraint (9) does not include the rental cost (since the household owns in the current period), but includes net income from selling the house. The costs are the current maintenance cost (δ), the selling cost (κ), and capital gains tax $q(s, h, p_1)$, which a homeowner has to pay above exemption level \bar{q} , on capital gains relative to the initial house purchase price \bar{p}_1 . Third, the interest rate is r for a net saver, but $r + \xi$ for a net borrower. Fourth, the consumption floor does not figure in rental cost. The household then begins the next period as a renter ($h' = 0$).

The problem of the homeowner who decides to stay in his house is characterized by:

$$V_1(i, s, b, m, x, p, h, a) = \max_{a' \geq -(1 - \lambda_i)hp_1} \left\{ u(c, h, s, 1) + \beta \sum_{s'} \pi_{s,s'}^s \sum_{m' > 0} \pi_{i,m,m'}^m \sum_{x'} \pi_{i+1,m',b,s',x'}^x V(i + 1, s', b, m', x', p', h, a') + \beta \pi_{m,0}^m v(hp'_1 + a') \right\} \quad (13)$$

subject to equations (6), (12) and:

$$c + a' + x + \delta hp_1 = (1 + \tilde{r})a + \psi_s b \quad (14)$$

$$c = \begin{cases} \max\{s\underline{c}, \tilde{c}\} & \text{if } h \leq \bar{h}p_1 \text{ and } a' \leq \min(a, 0) \\ \tilde{c} & \text{otherwise} \end{cases} \quad (15)$$

First, since a homeowner can borrow against the house, a' is not constrained from below by zero, but by $-(1 - \lambda_i)hp_1$. Second, in case the household does not survive to the next period, the estate includes the value of housing (hp'_1) as well as the financial asset position (a'). There is no capital gains taxation on the bequest of a house; by estate taxation laws in the U.S., the

Table 1: First-Step Estimated Parameters

Parameter	Description	Value ¹
μ_2	Household equivalence scale for 2-adult households	1.340
ψ_2	Income multiplier for 2-adult households	1.500
δ	Maintenance cost of housing	0.017
κ	House selling cost	0.066
r	Saving interest rate	0.040
ξ	Mortgage interest premium	0.016
g_0	Rental rate growth rate	0.011
g_1	House price growth rate	0.045
τ_q	Capital gains tax rate	0.190
\bar{q}	Capital gains tax deduction ²	250,000
\bar{p}_1	Purchase price of home (relative to 1996 price)	0.910
\bar{h}	Homestead exemption for consumption floor	357,000

¹ Annualized value.

² 1996 dollars.

exemption level for estate taxes were \$600,000 in 1996, and \$2,000,000 in 2006, so the tax does will apply for any housing bins in our model. For the same reason, we do not model taxation of financial asset bequests. Finally, the homeowner can access the social insurance program subject to the homestead exemption of $\bar{h}p_1$. The condition (15) also states that if the homeowner is in debt, the debt is not written off when the consumption floor is received.

5 Estimation

5.1 Estimation Strategy

Following Gourinchas and Parker (2002) and De Nardi et al. (2010), we use a two-step estimation strategy. In the first step, we estimate the parameters that are directly observable. Parameters associated with all the shocks and prices, as well as the initial conditions, are in this category. In the second step, given these exogenous parameters, we estimate the remaining parameters using a minimum-distance estimator, taking as targets the set of life-cycle profiles described above.

5.2 First Step Estimation

To match the HRS, we set the model period to two years. Each household can live up to 99 years of age, but there is a probability of an earlier death. We look at three cohorts corresponding to ages 63-67, 73-77, and 83-87 in 1996. We call them cohorts 1, 2, and 3, respectively, and define their age as the center point of each age bin (thus ages 65, 75, 85). For each cohort, we have

Table 2: Income Levels¹

	Cohort 1 (age 65)	Cohort 2 (age 75)	Cohort 3 (age 85)
Group 1	5831	6199	5520
Group 2	12049	9977	8055
Group 3	17844	13593	10481
Group 4	25868	18173	13743
Group 5	50227	37869	26090

¹ Annualized income in 1996 dollars.

six data observations that correspond to years 1996, 1998, 2000, 2002, 2004, and 2006. All the values that follow in this section are in 1996 dollars. Individual parameters from the first-step estimation are summarized in table 1.

Preferences. There is a variety of estimates for the household equivalence scale. We use $\mu_2 = 1.34$ for a two-adult household, following Fernández-Villaverde and Krueger (2007). This value is the mean of the existing estimates in the literature, ranging from 1.06 to 1.7.

Nonfinancial Income. In each cohort, we sort the households according to their nonfinancial income in 1996 and classify them into five bins, so that each bin carries approximately one-fifth of the total sample weight in 1996. The value of each bin is the average income of that bin’s households. Prior to binning, we adjust the income of two-adult households to make it comparable to that of singles. To do this, we use the sample of households who change size from two to one while in the sample, and compute the ratio of income when the household was a couple over the income after it became a one-adult household. The median of this ratio is 1.5, so we set the parameter $\psi_2 = 1.5$. That is, in the median, a two-person household that loses a spouse also loses about one-third of its income. Thus, we divide nonfinancial income of two-adult households by ψ_2 . Table 2 summarizes the resulting bins.

Household Size. Figure 12 presents the proportion of two-adult households conditional on age. The proportion is approximately linearly decreasing with age. For all the shock processes, we assume that the household size transition probabilities are time-invariant and estimate them from a pooled sample of all six waves of the HRS. Recall also that we abstract from remarriages, based on low occurrence of these among retirees in the data. Finally, we assume that all the transitions from two- to one-adult households are caused by death of the spouse, i.e. are involuntary. That is, we assume away divorce, which appears to be rare in our data.




Figure 12: Proportion of Two-Adult Households.

Table 3: Health Status Transition

Health status transition (age 65)					Health status transition (age 75)				
	Dead	Excellent	Good	Poor		Dead	Excellent	Good	Poor
Excellent	1.3	72.8	21.5	4.4	Excellent	3.9	60.1	26.9	9.2
Good	2.2	25.8	53.3	18.7	Good	6.6	21.1	46.9	25.4
Poor	9.6	6.1	20.7	63.7	Poor	16.3	3.8	17.6	62.3
Health status transition (age 85)					Health status transition (age 95)				
	Dead	Excellent	Good	Poor		Dead	Excellent	Good	Poor
Excellent	10.5	46.8	27.1	15.6	Excellent	28.5	29.5	19.8	22.3
Good	14.7	17.0	37.8	30.5	Good	32.9	12.9	26.8	27.5
Poor	28.8	5.1	13.2	52.9	Poor	56.9	4.2	13.6	25.3

Health Status and Mortality Shock Process. We group the five self-reported health categories in the HRS (excellent, very good, good, fair, poor) into three groups, combining the top two and the bottom two. We compute the probability that a respondent of health $m \in \{1, 2, 3\}$ is of health m' two years later, conditional on the respondent's age. In this procedure, we also compute the probability of death ($m' = 0$). Table 3 presents health transition probabilities for ages 65, 75, 85, and 95. As expected, first, the probability of dying is generally higher for older respondents, and those in worse health. Second, health status is persistent. Third, however, this persistence becomes weaker with age, as probability of death increases.

Medical Expenditures. As we mentioned before, we incorporate into our data the HRS exit waves, which collect medical expenditure information up to the end of life on respondents who die between two waves of the survey. We do not include expenses covered by Medicare, Medicaid or private health insurance since the corresponding model shocks are uninsurable. We estimate out-of-pocket (OOP) medical expenditure shocks by regressing the probability of zero OOP




Figure 13: Expected Medical Expenses of Singles

medical expenses, the mean and the standard deviation of log-medical expenses on household size, income, health status, a quartic in age, as well as interaction terms of age and the other variables. Under the assumption of log-normality, we construct expected OOP medical expenses from this estimation. In figure 13, we plot expected medical expenses for single households, by health status and by income. Medical expenditures increase with age, driven both by a rising mean and rising variability over time; the rise is particularly dramatic for those in poor health.

We then discretize the log-normal distribution using five grid points: zero, the mean, mean plus-minus one log standard deviation, and mean plus two times the log standard deviation. The last grid point captures the right tail of the distribution, emphasized by French and Jones (2004).

The estimates that we obtain imply that a 95-year-old single household of median income in bad health has a 5.5% probability over two years to get \$172,162 in OOP medical expenses, in 1996 dollars. A 95-year-old single household of high income in bad health has a 5.9% chance over two years to get \$318,037 in medical expenses. On average, about 2.9% of 95-year-old households in our estimation have more than \$200,000 of out-of-pocket medical expenses. These numbers are slightly below those computed by Ameriks et al. (2011) based on MEPS data, if we convert our numbers to 2006 dollars, as in their paper.

Housing. We approximate the distribution of house sizes in each cohort using ten grid points. We create this grid by classifying the households in each cohort into bins, each with about 10% of the sample, and using the mean house value within each bin as the grid points. Table 4 summarizes the result. In the model, we also restrict the choice of property values for renters to the same house bins for each cohort.

Table 4: House Size Distribution¹

	Cohort 1 (age 65)	Cohort 2 (age 75)	Cohort 3 (age 85)
Bin 1	21792	18267	16955
Bin 2	44935	37938	35743
Bin 3	63613	50803	47027
Bin 4	77839	64390	51910
Bin 5	88087	77583	62395
Bin 6	101358	88868	75851
Bin 7	125114	103150	88729
Bin 8	152107	137364	108380
Bin 9	195244	183191	148655
Bin 10	360683	345206	266577

¹ Value in 1996 dollars.

We set maintenance cost δ at 3.4% per two years (annually 1.7%). This is the value calibrated by Nakajima (2010) using data on depreciation of residential capital in National Income and Product Accounts. The selling cost of a house (κ) is set at 6.6% of the value of the house. This is the estimate obtained by Greenspan and Kennedy (2007). Grueber and Martin (2003) report the median selling cost of 7.0% of the value of the house.

Saving and Home Equity Borrowing. The saving interest rate is set at 8% (annually 4%). The mortgage debt premium ξ is set at 3.2% (annually 1.6%), which is the average spread between 30-year mortgages and Treasury bonds of the same maturity between 1977 and 2009.

Housing Prices. For the rate of growth of rents, g_0 , we use the CPI rent of primary residence component. The rate of growth for 1996-2006 averages at 2.1% per two years. We measure the rate of growth of house prices, g_1 , from the house price index (HPI) compiled by the Federal Housing Finance Agency. At 4.5% per year, this rate is just below the price change of 4.8% per year that we compute from self-reported home values in the HRS.

Policy Parameters. The last four parameters in table 1 refer to tax policy and social programs. We set the capital gains tax rate τ_q at 19%, which is the weighted average of marginal tax rates for the 1996-2006 period, during which the legislation changed repeatedly. The capital gains tax deduction \bar{q} is \$250,000, which is similarly averaged over this period. In order to calculate the capital gain, we assume that the house was initially purchased at $\bar{p}_1=91\%$ of the current home value, which is the average in our HRS sample; keeping track of initial house prices at household level is too costly computationally. Finally, we calibrate the homestead exemption level \bar{h} for the

Table 5: Initial Distribution – Selected Characteristics

	Cohort 1 (age 65)	Cohort 2 (age 75)	Cohort 3 (age 85)
Household size			
one-adult	0.28	0.56	0.76
two-adult	0.72	0.44	0.24
Health status			
1 (excellent)	0.50	0.39	0.33
2 (good)	0.27	0.32	0.28
3 (poor)	0.23	0.29	0.39
Tenure			
Homeowner	0.89	0.78	0.60
Renter	0.11	0.22	0.40
Net financial asset position			
Saver	0.82	0.92	0.96
Borrower	0.18	0.08	0.04


consumption floor based on De Nardi et al. (2016), who describe Medicaid rules state by state in 2009. We average homestead exemption levels and convert to 1996 dollars using the HPI.

Initial Distribution We construct the model’s initial distribution along the eight-variable state space from the 1996 HRS sample, simulate the model starting from this distribution, and use the outcome of the simulation to estimate the structural parameters in the second estimation step below. Table 5 shows the aspects of the initial distribution that we did not already describe.

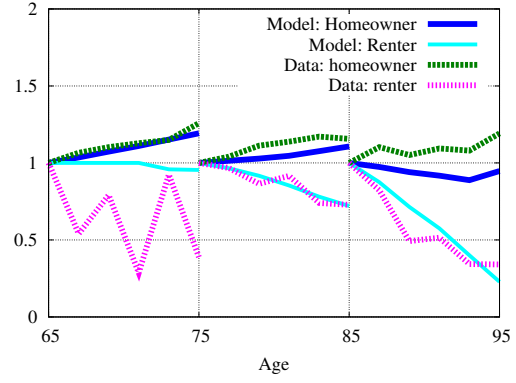
5.3 Second Step Estimation

In the second-step estimation, we choose parameters to fit the life-cycle profiles in section 3, using a minimum-distance estimator. The targets are cohort profiles of the homeownership rate, median total, financial and housing assets, proportion of households in debt, and median debt of debtors. We also target the normalized median net worth profiles for homeowners and renters, as in figure 2, as well as net worth by income bin and Medicaid participation rates by age.


Figure 14 gives the asset profiles of the three cohorts, comparing the model and data. The first panel presents median total asset profiles of each cohort, which is the classic statement of the retirement saving puzzle; the model matches these profiles perfectly. The second panel compares owners and renters in the model to those groups in the data, in terms of normalized median total asset profiles. For the first two cohorts, the model matches the data profiles fairly well, with the exception of the youngest cohort’s renters – but this group is small in the data




(a) Median Total Assets




(b) Owners vs Renters




(c) Total Assets by Income Bin



(d) Homeownership Rate



(e) Median Housing Assets



(f) Median Financial Assets

Figure 14: Benchmark Model Fit - Asset Holding Profiles

and the profile is noisy. For the oldest cohort, the model underpredicts somewhat the extent of asset accumulation of owners, but generates close to the full empirical difference in the behavior of the two groups. The third panel shows total assets by income group. The model generally does well on these. For the youngest cohort, the model underpredicts net worth accumulation




Figure 15: Benchmark Model Fit - Debt and Medicaid Profiles

of the top tail, while it overpredicts it somewhat for the oldest cohort. The remaining panels of the figure show that the model nearly perfectly replicates homeownership rates, median housing assets of homeowners and median financial assets in the model relative to data.

The top two panels of figure 15 present the proportion of households in debt (panel (a)) and median debt among indebted households (panel (b)). Given our parsimonious modeling of equity borrowing costs, we cannot match both the extensive and intensive margins of debt. The model matches the rate of indebtedness very well, but does less well on debt levels. In the data, the few households that do continue to borrow late in life appear to borrow quite a bit; in our model, the estimated tight collateral constraints make that impossible. Matching the intensive margin would require a more flexible model of mortgage contracts (see Nakajima and Telyukova (2017)), at the cost of additional computational burden. Here, the intensive-margin discrepancy is acceptable as it concerns a small subset of the population.

The bottom two panels of figure 15 show the rate of Medicaid reciprocity for the sample as

Table 6: Second-Step Parameter Estimates

Parameter	Description	Value
β	Discount factor ¹	0.96
η	Consumption aggregator	0.81
σ	Coefficient of RRA	2.93
ω_1	Extra-utility from ownership	2.52
γ	Strength of bequest motive	7.19
ζ	Shifter of bequest utility	45714
\underline{c}	Consumption floor per adult ¹	8981
λ_{65}	Collateral constraint for age-65	0.64
λ_{75}	Collateral constraint for age-75	0.96
λ_{85}	Collateral constraint for age-85	0.99
λ_{99}	Collateral constraint for age-99	1.00

¹ Biennial value.

a whole, and by homeownership status. For the sample as a whole, the model does quite well, slightly underpredicting the rate for the second cohort, and overpredicting it for the third. From the second panel, it is clear that the discrepancy in Medicaid rates is due to renters rather than owners in the model. This is important, because in the model, as in the data, there is a homestead exemption for Medicaid eligibility that could motivate some homeowners to stay in the house longer. The graph shows that we do not overstate the influence of Medicaid on homeownership. The fact that the model underpredicts the share of renters on Medicaid is explained by the fact that we only have five income bins in the model; by design, we cannot replicate the lowest tail of the income distribution, a key source of Medicaid claims in the data.

The resulting parameters from the second-step estimation are in table 6. First, the estimated collateral constraints imply that upon retirement, households' ability to borrow against home equity quickly becomes limited. Retirees of age 65 can borrow up to 36% of their home equity. The constraint tightens considerably by the time they are 75, to just 4% of home equity, and to just 1% for homeowners who are 85 years old. These numbers should not be interpreted literally as downpayment constraints; rather, these constraints capture the overall cost of equity borrowing, which reflect, for example, the fact we mentioned before that many retirees fail the income requirement for equity loans (Caplin (2002)), so that after retirement, equity borrowing becomes much costlier.⁷

Based on the Medicaid reciprocity rates, we estimate the consumption floor to be about \$4,490

⁷A related matter is that reverse mortgages are instruments available specifically to the elderly, but in separate work, we find these to be very costly, due to insurance costs, which is consistent with our interpretation of λ_i as capturing high costs of borrowing in retirement. See Nakajima and Telyukova (2017) for details.

per person per year in 1996 dollars. Our estimate lines up well with Hubbard et al. (1994), who measure the non-Social-Security consumption floor for the elderly to be \$6,893 per *household* per year in 1984 dollars. Our estimated parameter of extra utility from homeownership is 2.5, which means that it is 3.5 times more appealing to own a home than it is to rent. As we mentioned before, this parameter captures all possible nonfinancial benefits of homeownership, as well as those financial benefits that are not explicitly in the model. The estimate that we have is consistent also with the findings in Venti and Wise (2004), who find strong support in the data for an AARP survey statement that retirees like to stay in their current residence as long as possible.

The relative risk aversion parameter that we estimate is $\sigma = 2.93$, in the middle of the range used in consumption-saving literature. Thus, we do not require a high degree of risk aversion to match the slow rate of asset decumulation that we observe in the data. The estimated strength of the bequest motive is 7.19. The interpretation of this and remaining parameters will be clear in the next sections in the context of our experiments.

6 Experiments: Decomposing the Retirement Saving Puzzle

We now use the estimated model to evaluate the quantitative contribution of key model features to retiree saving behavior, by shutting down these mechanisms one at a time, keeping all other model features constant, and comparing the results to the benchmark model outcome. At the end of this section we quantify the role of all the model features, by stripping the model down to the basic life-cycle case with only longevity risk, then re-introducing the features back in different orders. The housing mechanisms that we focus on are extra utility of homeownership, collateral constraints, and the housing boom of 1996-2006. We also evaluate the impact of medical expenses and bequest motives, to follow previous literature. We focus on the most salient results for each experiment; the full set of results is available in the online appendix.⁸

6.1 Role of Extra Utility of Homeownership

To evaluate the role of nonfinancial and financial benefits of homeownership in retiree saving decisions, we set $\omega_1 = \omega_0 = 1$, so that owned and rented homes are the same in terms of utility. Panel (c) of figure 16 shows that, not surprisingly, the utility benefit of homeownership encourages retirees to own homes, to a similar extent across cohorts, and from panel (d) it is also clear that it suppresses home equity borrowing, particularly for the youngest cohort. Through extra homeownership, the utility benefit impacts net worth somewhat, as panel (a) demonstrates;

⁸The online appendix can be found at <http://dss.ucsd.edu/itelyuko/research.html>.

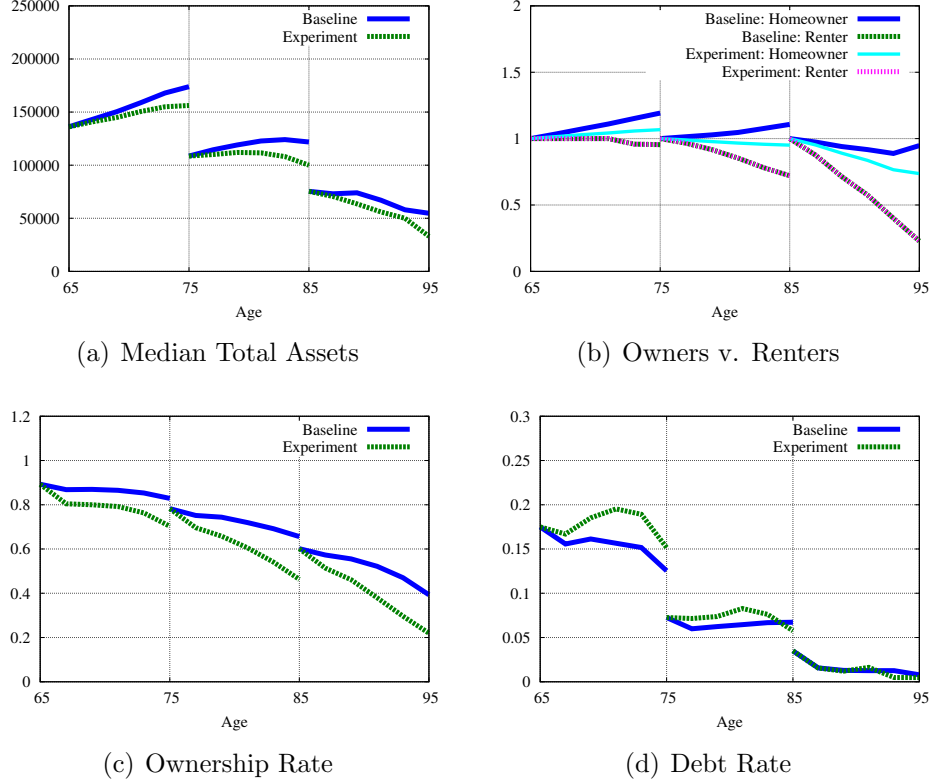


Figure 16: Decomposition Experiment – No Extra Ownership Utility

overall, however, the contribution of this channel to the retirement saving puzzle for the median household is mild. The extra utility of ownership affects only homeowners (panel (b)), and the difference in the rate of dissaving between owners and renters remains large.

6.2 Role of Collateral Constraints

For this experiment, we relax the collateral constraints on homeowners, make them uniform across age. For this purpose, we set everyone's collateral constraint λ to 0.2, allowing all homeowners to borrow up to 80% of their equity, which is a standard feature of most mortgage contracts.

As figure 17 shows, the illiquidity of housing is sufficient to create most of the observed decline in the homeownership rate, all else equal (panel (c)). That is, homeowners in our model sell their homes largely when they are forced against their borrowing constraint. The borrowing constraints depress borrowing (panel (d));⁹ if homeowners were able to access their equity freely, they would hold on to their house, motivated by the utility benefit and housing boom, and

⁹For ease of comparison, we keep the scale of the graphs constant, at the cost of having this graph partially off the scale.

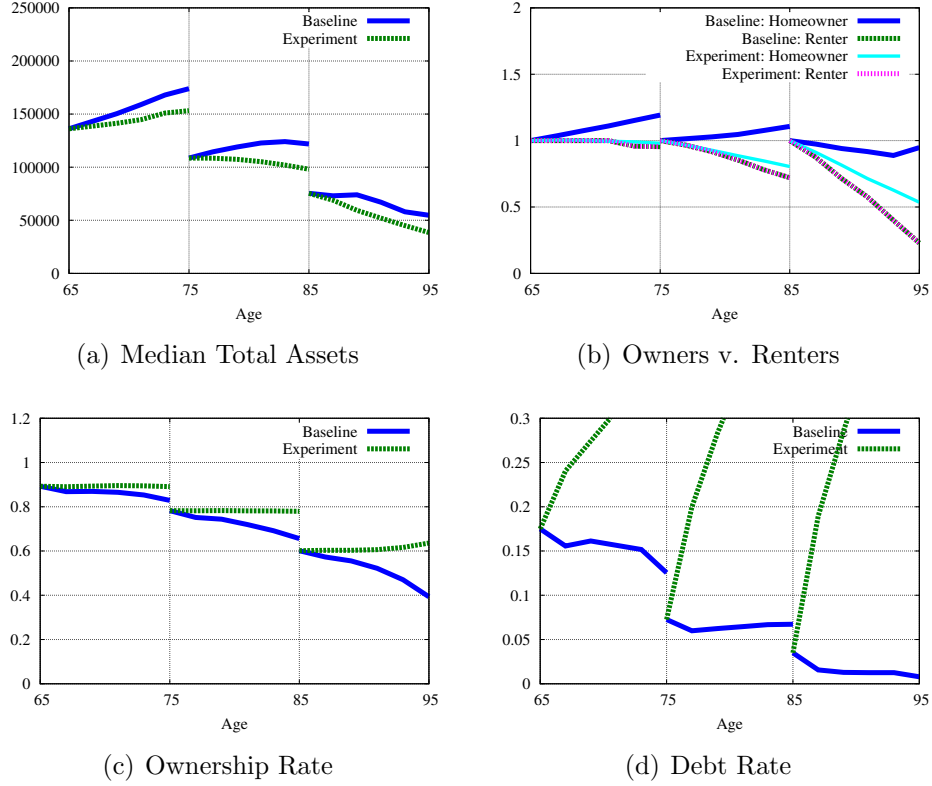


Figure 17: Decomposition Experiment – Uniform Collateral Constraints

borrow against their equity instead, decumulating their housing asset more quickly. This in turn is reflected visibly in the median net worth in panel (a), suggesting that collateral constraints contribute to the retirement saving puzzle in a significant way, especially for younger cohorts.

Finally, the collateral constraint, and thus illiquidity of housing, contributes significantly to the difference in dissaving rates between homeowners and renters (panel (b)). Only for the oldest cohort there remains a noticeable difference, even with the relaxed collateral constraint.

6.3 Role of the Housing Boom

To evaluate the role of the 1996-2006 housing boom, we shut down the exogenous increasing trend to the aggregate housing price, making it constant instead. The obvious impact is the flattening of median net worth profiles in panel (a) of figure 18, which happens through two channels. First, there is a significant flattening of median housing asset profiles. In addition, panel (d) shows a decline in the debt rate when there is no housing boom; the median amount of debt, especially in the younger cohort, also declines somewhat. (See online appendix). This is intuitive: a housing boom increases the amount of home equity that households can extract,

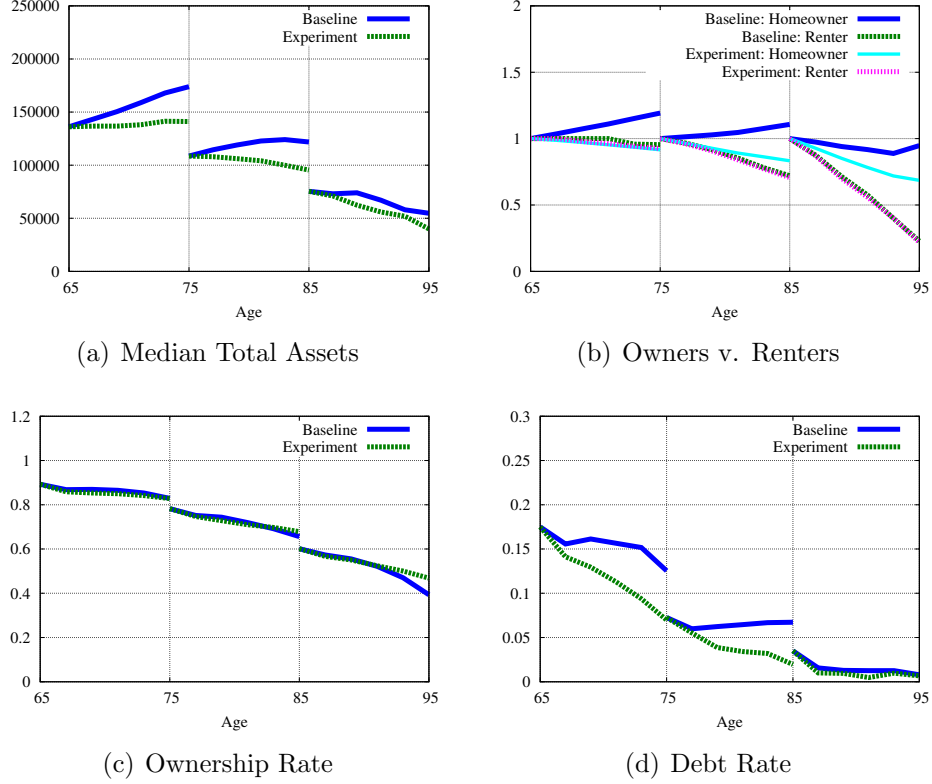


Figure 18: Decomposition Experiment – No Housing Boom

and some households take advantage of that. The overall impact of the housing boom is an increase in the median net worth for all cohorts; this impact is least pronounced for the oldest cohort. Note also that the housing boom, like collateral constraints, plays an important role in creating the differences in the dissaving behavior of homeowners versus renters (panel (b)), with the oldest cohort being the biggest exception.

6.4 Role of Medical Expenses

Following De Nardi et al. (2010), in this section we perform two experiments. First, we shut down medical expense risk, setting everyone's expenses to the mean, as in figure 13. Then, we set everyone's medical expenses to $x = 0$. The two experiments have similar effects, so we show only the risk experiment here. First, we see that homeownership is not affected by medical expense risk (panel (c), figure19), except for the oldest cohort, a small fraction of which would sell their homes if we shut down medical expense risk. This effect comes from the homestead exemption of Medicaid: in the presence of medical expense risk, some of the oldest homeowners have an incentive to hold on to their home and use Medicaid if a big medical expense shock realizes, rather than to sell the home and use home equity to pay for the expense. Through this channel,

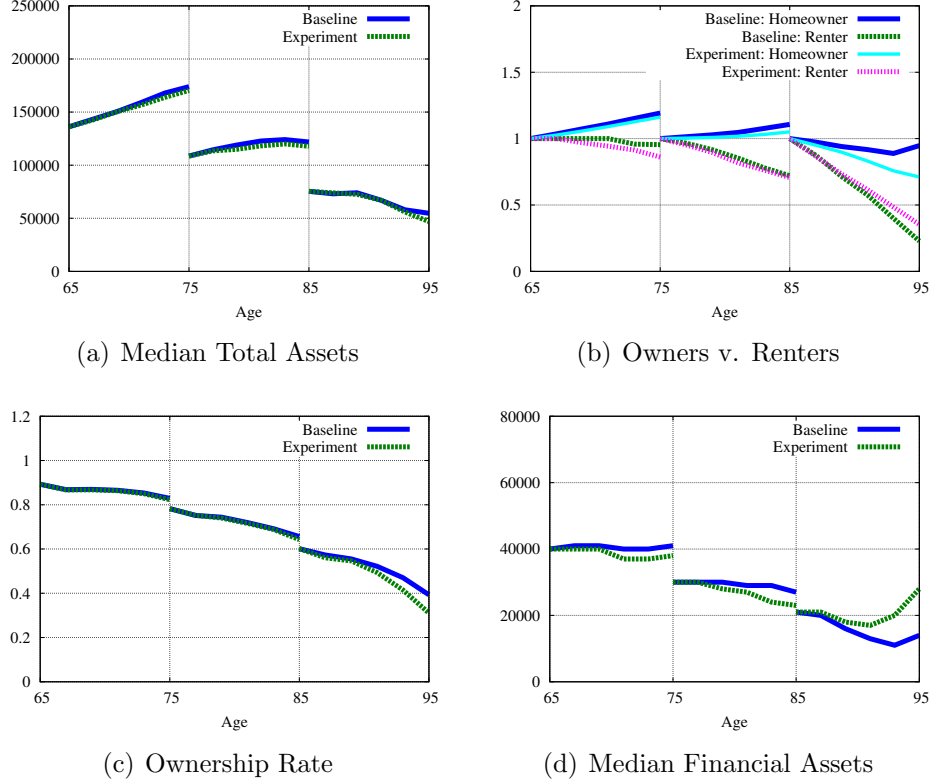


Figure 19: Decomposition Experiment – No Medical Expense Risk

medical expenses contribute slightly to the difference between homeowners and renters (panel (b)) in the oldest cohort. This is confirmed if we look at housing assets by income bin (see online appendix): the housing assets of the oldest and poorest retirees are most affected.

The overall effect on median net worth is small (panel (a)). In the absence of medical expense risk, the two younger cohorts decumulate financial assets more rapidly (panel (d)) pointing to a moderate precautionary motive for saving, while median housing assets are not affected. For the oldest cohort, financial assets increase in the absence of expense risk: this is due to those who would sell the house in this case, thus liquidating and spending down some of their home equity, but saving the rest in financial assets. The overall impact is a slight decline in total assets.

6.5 Role of Bequest Motive

Figure 20 shows that if we shut down the bequest motive by setting $\gamma = 0$, there is a noticeable change in median net worth. The bequest motive increases the homeownership rate for the middle and especially the last cohort somewhat, and discourages equity borrowing, as panel (d) shows. As we would expect, it impacts both homeowners and renters, although the difference

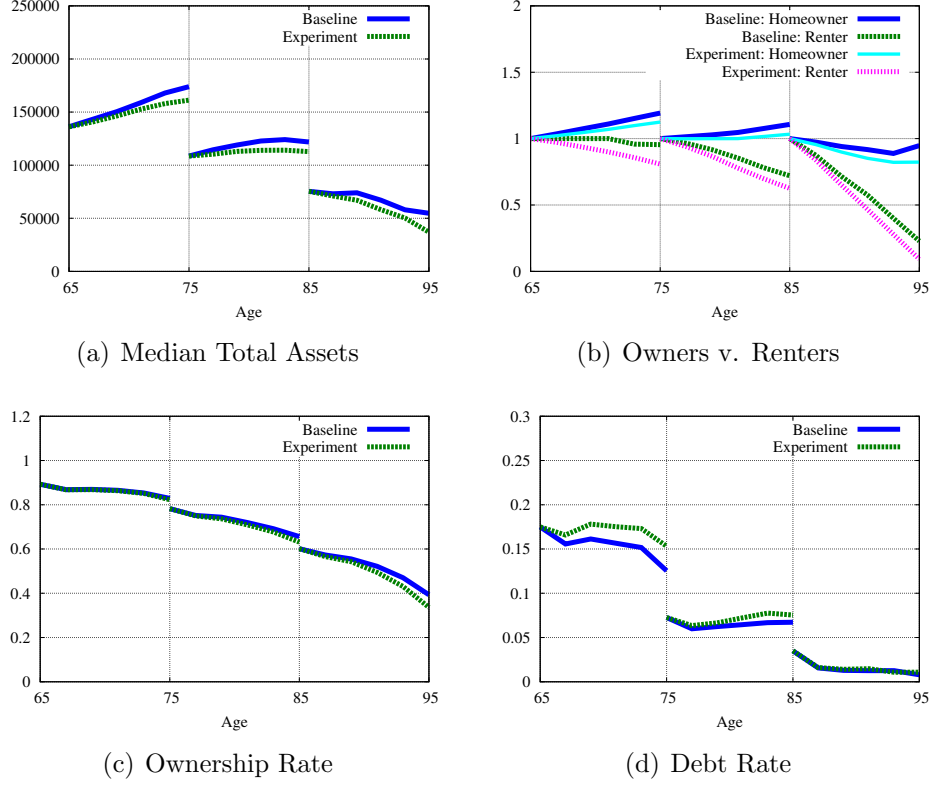


Figure 20: Decomposition Experiment – No Bequest Motive

between homeowners and renters remains significant even without the bequest motive.

6.6 Summary: Individual Decomposition Results

In the data, many households are homeowners well into their old age, and this fact turns out to be crucial for the retirement saving puzzle. As the experiments above show, retired homeowners choose to remain owners late in life due to a combination of financial and nonfinancial benefits of ownership (expressed in our utility parameter ω_1 , avoidance of selling costs, capital gains taxes and the like), as well as the expected house price growth. Once retirees choose to stay in their homes, they become beneficiaries of the housing boom, which significantly contributed to the retirement saving puzzle over the period 1996-2006. Finally, as homeowners age, they become increasingly locked into their home equity because of tightening borrowing constraints; this also contributes to the flatness of the net worth profiles that we observe. On the flip side, retirees can hit their borrowing constraints – due to medical and other expenses – and can be forced to sell the house as a result, becoming renters instead.

The benchmark model also suggests a moderate contribution of medical expense risk to the

retirement saving puzzle, and a somewhat larger impact of bequest motives. Medical expenses do not interact with homeownership except for the oldest old through the Medicaid channel; the bequest motive interacts more noticeably, as one reason for homeownership late in life.

The decompositions above imply possible further interactions between the model channels that we are concerned with. Bequest motives and utility benefits of homeownership may interact and reinforce each other in creating motives for homeownership; medical expenses may interact with borrowing constraints and cause some homeowners to sell. In order to tie the experiments together and understand the quantitative contribution of each channel, we next strip the model down to a basic lifecycle framework with longevity risk, and then re-introduce those channels in succession. We do this in several different orders, to highlight the interactions of interest.

6.7 Quantitative Assessment of the Retirement Saving Puzzle

For these sequential experiments, we focus on the most standard statement of the retirement saving puzzle – the median net worth of each age group. Taking the benchmark net worth median as 100%, we quantify how much of it, relative to the benchmark, is generated for each cohort by each model mechanism. We present results for selected age groups in table 7.

In the stripped-down model, the only source of saving is precautionary motive against longevity risk. In this model, the median age-75 household would have 50% of that age group’s empirical and benchmark-model net worth; at age 85, the net worth would be at 44% of the observed amount; and by age 95, the median net worth is at 33% of the benchmark. Households of age 99 in the model have zero wealth, because they die with certainty after that period.

Each numbered panel of table 7 represents a different order in which we re-introduce the model features. We vary the order of bequest motives with medical expense risk, and with the three housing-related channels. In addition, we investigate changes in the ordering of the housing-related channels themselves, first introducing the housing boom before the collateral constraint, and then after. The overview is that housing plays an important quantitative role for the retirement saving puzzle, accounting for between 28 and 53% of the observed median net worth, depending on age and order. In addition, there is visible interaction of the bequest motive with housing channels, because its quantitative contribution changes depending on the order in which it is introduced. The bequest motive accounts for 8-21% of median net worth of retirees, depending on age. Finally, the quantitative role of medical expenses is moderate, 0-14% depending on age, and there is moderate interaction with housing for the oldest cohort.

The role of the three housing channels together ranges between 28 and 40% for the youngest group, depending on whether their preferences have a bequest motive introduced first, and be-

Table 7: Contribution of Each Channel to Median Net Worth – Sequential Decompositions, Incremental Contributions

Model	Age 75 (cohort 1)	Age 85 (cohort 2)	Age 95 (cohort 3)
Simple life-cycle with longevity risk	50	44	33
1 Bequest motive	20.6	17.7	14.6
Utility benefit of ownership	5.2	4.2	5.0
Housing boom	11.0	15.2	18.3
Collateral constraints	11.6	15.8	15.0
Medical expenditure risk	2.0	3.3	14.2
<i>Total housing</i>	<i>27.8</i>	<i>35.2</i>	<i>38.3</i>
2 Utility benefit of ownership	4.0	3.7	1.3
Housing boom	14.4	15.0	21.9
Collateral constraints	21.6	19.6	9.6
Bequest motive	8.3	14.6	20.0
Medical expenditure risk	2.0	3.3	14.2
<i>Total housing</i>	<i>40.0</i>	<i>38.3</i>	<i>32.9</i>
3 Medical expenditure risk	7.8	8.8	-0.5
Bequest motive	14.4	12.2	15.1
Utility benefit of ownership	4.1	3.3	7.0
Housing boom	12.1	12.4	15.8
Collateral constraints	11.9	19.5	29.7
<i>Total housing</i>	<i>28.1</i>	<i>35.2</i>	<i>52.5</i>
4 Bequest motive	20.6	17.7	14.6
Medical expenditure risk	1.7	3.3	0.0
Utility benefit of ownership	4.1	3.3	7.0
Housing boom	12.1	12.4	15.8
Collateral constraints	11.9	19.5	29.7
<i>Total housing</i>	<i>28.1</i>	<i>35.2</i>	<i>52.5</i>
5 Bequest motive	20.6	17.7	14.6
Medical expenditure risk	1.7	3.3	0.0
Utility benefit of ownership	4.1	3.3	7.0
Collateral constraints	5.2	10.2	18.5
Housing boom	18.9	21.7	26.9
<i>Total housing</i>	<i>28.1</i>	<i>35.2</i>	<i>52.5</i>

tween 33 and 53% for the oldest households. Of these three channels, the utility benefit of homeownership is the least influential quantitatively, contributing 4-5% of median net worth of the youngest cohort, and 1-7% for the oldest cohort. Both the housing boom and the illiquidity of housing inherent in the collateral constraints are more significant contributors to net worth. Collateral constraints are responsible for 5-22% of net worth for the youngest cohort, and for 10-30% for the oldest cohort, depending on the experiment, while the housing boom contributes respectively 11-19% and 16-27% to the youngest and oldest cohorts' median net worth profiles.

The first experiment shows that adding the bequest motive to the model with only longevity risk (experiment 1) implies that it accounts for 21% of the median net worth of age-75 households, for 18% of the net worth of age-85 households, and 15% for households of age 95. If the bequest motive is instead introduced after the housing channels (experiment 2), its contribution declines for the youngest retirees, to 8% of observed median net worth, but increases for the oldest households, to 20%. The change in the numbers indicates that the bequest motive affects the behavior of homeowners, and the decision to own a home for the oldest cohort.

Medical expense risk, in the model with only longevity risk, accounts for 8% of savings of 75-year-olds, but for under 2% if they have a bequest motive (experiments 3 and 4). For the oldest households, the impact is near zero. The role of bequest motives remains significant and robust when introduced into the economy with medical expense risk. There is some interaction between the two motives, but it is not quantitatively significant. From experiments 1 or 2, we see that when medical expenses are introduced into the model with both housing and bequest motives, their role lies in the middle of this range for the younger cohorts, but it is amplified to 14% for the oldest households. This reflects the Medicaid homestead exemption, which allows some older households to retain a lot of illiquid wealth that they would otherwise spend down.

7 Identification and Sensitivity

While in this model, as in all models of its type, we are unable to provide a formal identification argument, and while calculation of standard errors is hindered by the computation time, in this section we discuss what estimation target determines each parameter, using sensitivity analysis. We also show that our key result, the quantitative role of housing, is robust to parameter perturbations. Finally, we show that the two parameters that determine bequest and precautionary motives are difficult to disentangle. Details of these experiments are in appendix Appendix B.

7.1 The Role of Housing and Parameter Identification

The role of housing in our model is determined by three channels which correspond to three parameters. The first is the collateral constraint, with the age-dependent parameter λ_i , which is identified by the debt rate profile. The second channel is the housing boom, which is driven entirely by the house price p with the growth rate g_1 , both of which are exogenous. These parameters guiding the contribution of the housing boom and the collateral constraints are identified nearly one-for-one by the respective estimation targets or exogenous assumption. The third channel is the utility benefit of homeownership, with parameter ω_1 , which helps pin down by the homeownership rate. Since ω_1 affects only homeowners and not renters, it also helps pin down the separation between the normalized net worth profiles of owners versus renters. We cannot claim one-for-one identification for ω_1 or the rest of the parameters. In a series of experiments here, we perturb one parameter at a time, holding all others fixed, to show how the model fit is affected, and how the role of housing responds to the change.

The consumption floor parameter \underline{c} determines the proportion of households who use Medicaid in the model. \underline{c} also helps generate the separation between homeowners' and renters' net worth profiles, since renters rely on the consumption floor more often, and thus are more sensitive to changes in \underline{c} . The benchmark estimate of \underline{c} is consistent with the level of social benefits estimated in other studies, and the benchmark model matches the overall Medicaid claim rate, but it underestimates the use of Medicaid among renters. Here, we re-calibrate \underline{c} to match the profile of Medicaid take-up rate among renters, at the expense of Medicaid take-up rates among owners, as well as the median financial asset and net worth profiles of renters (see appendix B.1). The re-calibration gives $\underline{c} = 11,600$; with this value, the three housing channels still account for 29-41% of median net worth at age 75, and for 53-63% at 95.

The value of ω_1 , as discussed above, is affected by two estimation targets. Here we re-calibrate ω_1 to correct the small discrepancy in homeownership rates between the benchmark model and the data; this comes at a cost of a small deterioration in the normalized net worth profiles of homeowners. With the result of $\omega_1 = 1.5$, housing accounts for 26-40% of median net worth at age 75, and 31-47% of net worth at age 95.

The aggregator of housing and non-housing consumption in the utility function, η , changes the balance between housing and financial assets. In particular, a lower η yields a higher homeownership rate, a higher total asset profile of homeowners, and a lower financial asset profile. Here, we recalibrate η to better match the homeownership rate, at the expense of financial asset and net worth profiles, and the Medicaid rates among owners, which decline. With $\eta = 0.86$, housing contributes 25-38% of median net worth at age 75, and 26-42% of net worth at age 95.

The discount factor β helps replicate the total and financial asset holdings, especially of younger cohorts. We test a re-calibration of β that better matches the median financial asset profile, at the expense of the overall net worth profile, and of the renters' net worth profile. This yields $\beta = 0.99$. The role of housing remains robust to this change, at 25-36% of median net worth at age 75, and 33-48% at age 95.

Finally, the bequest shifter ζ affects the difference in bequest motives between wealthy and less wealthy households. We re-calibrate ζ to match the median financial profile better, with the resulting value of 100,000. This leads to a significant deterioration in the fit of the homeownership rate for the oldest cohort, of the debt rate and of the separation of normalized net worth between borrowers and renters. With this, housing contributes 24-38% of the youngest cohort's median net worth, and 22-44% of the oldest cohort's.

In sum, these experiments give a sense that each parameter plays a distinct role in fitting the benchmark model to all of our estimation targets, and the result of each perturbation is a noticeable deterioration of fit. It is also encouraging that the role of housing remains robust to the perturbation of the value of these parameters.

7.2 Strength of Bequest and Precautionary Motives

The second set of our benchmark results concerns the role of bequest motives and medical expenses, two contributors to the retirement saving puzzle emphasized in previous literature, in generating the observed median net worth profiles. This role is primarily determined by the parameters σ , the curvature of the utility function, and γ , the strength of the bequest motive. It is difficult to make a formal argument about identification of these two parameters: while our experiments suggest that we may be able to use age and wealth heterogeneity to identify them, because bequest motives and medical expenses affect the young and old, rich and poor differently, it is also clear that these parameters interact. This is an issue that would affect any life-cycle model of the kind we use here.

In order to demonstrate the sensitivity of our results to the parameters σ and γ , here we perform two extreme experiments. First, we set $\sigma = 3.8$, as in De Nardi et al. (2010), a value that is nearly 25% above our estimated value. Then, holding all other parameters constant, we re-estimate γ such that the median net worth profile is matched. In this case, we find that $\gamma = 0$ is required, and this model cannot match all of the other targets equally well (appendix B.2). Specifically, the fit of median financial asset profiles deteriorates relative to the benchmark, especially for the top income quintile of the youngest and oldest cohorts, so that this version of the model cannot be considered on par with the benchmark. Also, in this specification, the

renters, particularly in the middle cohort, do not decumulate assets quickly enough relative to the data. Nevertheless, this altered model still predicts a robust and strong role for housing, which accounts for 35-38% of median net worth profiles of the youngest cohort, and 32-53% for the oldest cohort. In light of the above discussion, this is not surprising: the relevant parameters do not significantly interact with σ and γ . However, now we observe that the quantitative role of bequests declines to zero, while the role of medical expenses rises, to 8-11% for age-75 retirees, and -0.6 - 22% for the older retirees.

In the second experiment, we again raise σ to 3.8, but now we keep γ at its benchmark estimated value, and instead let β adjust to match the median net worth profile. In this case, $\beta = 0.929$ and the model deteriorates along other dimensions: the median financial asset profile decreases relative to the benchmark and data, while the debt rate rises, for the youngest cohort. This makes intuitive sense: less-patient households tend to save less. However, this version of the model again keeps the role of housing robust, at 32-41% and 28-49%, respectively, for age-75 and age-95 retirees, medical expenses still play a somewhat expanded role relative to the benchmark case, at 5-11% and 0.2-22% respectively, but bequest motives are also present, contributing 4-13% to younger retirees' median net worth, and 5-9% for oldest retirees.

The message of these experiments is, first, that our main result regarding the role of housing is robust to the changes of parameters σ , γ and β in the model. However, second, in life-cycle models of retirement with warm-glow bequest motives, it may be difficult to disentangle the strength of bequest motives from precautionary motives. This point was also made by Ameriks et al. (2011). Thus, future research will need additional strategies for identification, which may involve, for example, more detailed modeling of bequests, or the use of data from multiple countries, or both. Nakajima and Telyukova (2018) take a step in that direction.

8 Experiment: The Role of Home Maintenance

In this section, we conduct one more experiment to evaluate the role of home maintenance as a possible “hidden” channel of asset decumulation. We are motivated by the study of Davidoff (2006), who finds that elderly homeowners spend on average 0.8% less per year on home maintenance than younger owners of a similar house, and that similar houses sell at lower prices if the owner was over 75 years old.

In order to conduct this experiment, we add a choice margin to our model. Whereas in the benchmark, we assume that all households pay the maintenance cost (1.7% of equity per year) and as a result, the house does not depreciate, now we give the households a choice: they can continue to pay the cost, or they can choose not to pay it, but the house will depreciate as a result

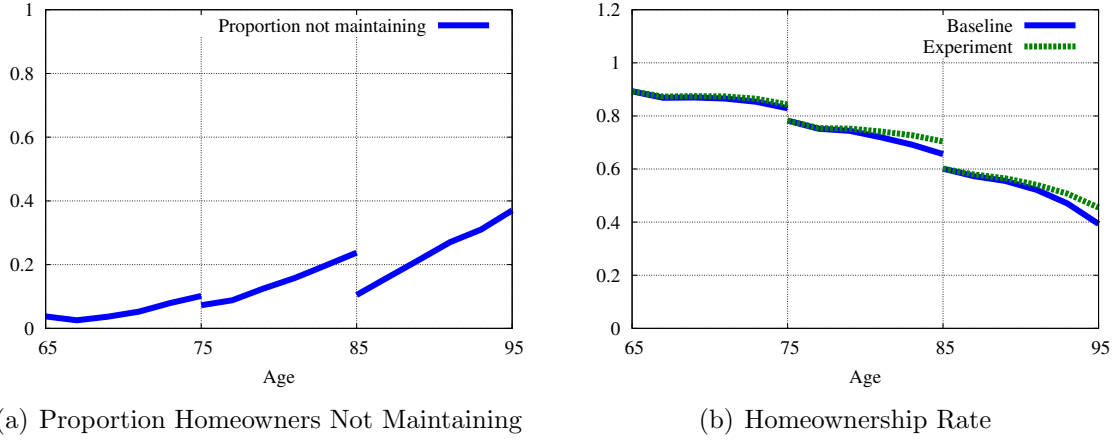


Figure 21: Experiment: Maintenance Decision

at the rate of 4.25% per year. The calibration for this experiment came from the maintenance spending number in Davidoff (2006) mentioned above. In addition, he calculates that houses of older homeowners appreciate at a rate of 2 percentage points per year lower than houses of younger owners. These two numbers together, in addition to the maintenance cost parameter, give us the parameters specified here.

An additional important point is that in order to conduct this experiment, we make the assumption that self-reported home values do not take into account depreciation of the house. That is, we assume that homeowners who stay in their homes do not have their house appraised and are not aware of the rate at which their home depreciates. This assumption is supported by Venti and Wise (2004), who find that self-reported home values are exaggerated. Still it may be an extreme assumption, and thus the results of this experiment should be treated with this caveat in mind. On the other hand, the empirical findings of Davidoff (2006) appear to be that the maintenance margin is an important one, and we can use our model to get a sense of how much it might contribute to the retirement saving puzzle in a hidden way.

Figure 21 gives the results of the experiment. First, the proportion of homeowners who choose not to maintain grows steadily with age (panel (a)), which is consistent with empirical findings in the literature. While in the youngest cohort, few homeowners choose the non-maintenance option, by age 95, 37% of homeowners opt not to maintain. Intuitively, the older a homeowner gets, the more likely she is to be borrowing-constrained. Choosing not to maintain gives her an alternative way to tap into some of her equity indirectly. This is mirrored in panel (b), where more older homeowners are able to keep their homes, compared to the benchmark model.

9 Conclusion

In this paper, we study homeownership in retirement, to understand what role it plays in accounting for the retirement saving puzzle. We do so by estimating a model of retiree saving that is explicit about differentiating between housing and nonhousing assets, targeting jointly in estimation not only median net worth lifecycle profiles of retirees, but also profiles of homeownership rates, housing and financial assets separately, renters' and homeowners' net worth separately, as well as debt rates and amounts. In our estimated model, housing plays a key role in accounting for the retirement saving puzzle, through a combination of utility benefits of homeownership, illiquidity of housing, and the housing boom of 1996-2006. Moreover, bequest motives play a significant role, while medical expenses play a moderate role once housing assets and homeownership are explicitly taken into account. We demonstrate also, however, that the relative roles of bequest motives and medical expenses are difficult to identify separately in this class of models. The quantitative role of housing is not significantly affected by these two channels, and is robust to perturbation of the other parameters. Relative to previous literature, conclusions regarding the retirement saving puzzle change if one considers housing and motives for homeownership late in life explicitly, and separately from overall net worth of retirees.

References

- Ameriks, John, Andrew Caplin, Steven Paufer, and Stijn van Nieuwerburgh**, “The Joy of Giving or Assisted Living? Using Strategic Surveys to Separate Public Care Aversion from Bequest Motives,” *Journal of Finance*, 2011, *66* (2), 519–561.
- Caplin, Andrew**, “Turning Assets into Cash: Problems and Prospects in the Reverse Mortgage Market,” in Olivia S. Mitchell, Zvi Bodie, Brett Hammond, and Steve Zeldes, eds., *Innovations in Retirement Financing*, Philadelphia, PA: University of Pennsylvania Press, 2002, chapter 3, pp. 234–253.
- Chambers, Matthew, Carlos Garriga, and Don E. Schlagenhauf**, “The Loan Structure and Housing Tenure Decisions in an Equilibrium Model of Mortgage Choice,” *Review of Economic Dynamics*, 2009, *12* (3), 444–468.
- Davidoff, Thomas**, “Maintenance and the Home Equity of the Elderly,” 2006. Unpublished.
- Davis, Morris A. and Jonathan Heathcote**, “Housing and the Business Cycles,” *International Economic Review*, 2005, *46* (3), 751–784.
- De Nardi, Mariacristina**, “Wealth Inequality and Intergenerational Links,” *Review of Economic Studies*, 2004, *71* (3), 743–768.
- , **Eric French, and John Bailey Jones**, “Why Do the Elderly Save? The Role of Medical Expenses,” *Journal of Political Economy*, 2010, *118* (1), 39–75.
- , —, and —, “Medicaid Insurance in Old Age,” *American Economic Review*, 2016, *106* (11), 3480–3520.
- Díaz, Antonia and María J. Luengo-Prado**, “The Wealth Distribution with Durable Goods,” *International Economic Review*, 2010, *51* (1), 143–170.
- Fernández-Villaverde, Jesús and Dirk Krueger**, “Consumption over the Life Cycle: Facts from Consumer Expenditure Survey Data,” *Review of Economics and Statistics*, 2007, *89* (3), 552–565.
- and —, “Consumption and Saving over the Life-Cycle: How Important are Consumer Durables?,” *Macroeconomic Dynamics*, 2011, *15* (5), 725–770.
- French, Eric and John Bailey Jones**, “On the Distribution and Dynamics of Health Care Costs,” *Journal of Applied Econometrics*, 2004, *19* (4), 705–721.

- Gervais, Martin**, “Housing Taxation and Capital Accumulation,” *Journal of Monetary Economics*, 2002, *49* (7), 1461–1489.
- Gourinchas, Pierre-Olivier and Jonathan A. Parker**, “Consumption over the Life-Cycle,” *Econometrica*, 2002, *70* (1), 47–89.
- Greenspan, Alan and James Kennedy**, “Sources and Use of Equity Extracted from Homes,” 2007. Federal Reserve Board, FEDS Working Paper No. 2007-20.
- Grueber, Joseph and Robert F. Martin**, “The Role of Durable Goods in the Distribution of Wealth: Does Housing Make Us Less Equal?,” 2003. Board of Governors of the Federal Reserve System, International Finance Discussion paper No. 773.
- Hubbard, R. Glenn, Jonathan Skinner, and Stephen P. Zeldes**, “The Importance of Precautionary Motives in Explaining Individual and Aggregate Saving,” *Carnegie-Rochester Conference Series on Public Policy*, 1994, pp. 59–125.
- , – , and – , “Precautionary Saving and Social Insurance,” *Journal of Political Economy*, 1995, *103* (3), 360–399.
- Hurd, Michael D.**, “Mortality Risk and Bequests,” *Econometrica*, 1989, *57* (4), 779–813.
- Kopecky, Karen A. and Tatyana Koreshkova**, “The Impact of Medical and Nursing Home Expenses on Savings,” *American Economic Journal: Macroeconomics*, 2014, *6* (3), 29–72.
- Li, Wenli and Rui Yao**, “The Life-Cycle Effects of House Price Changes,” *Journal of Money, Credit, and Banking*, 2007, *39* (6), 1375–1409.
- Lockwood, Lee M.**, “Bequest Motives and the Annuity Puzzle,” *Review of Economic Dynamics*, 2012, *15* (2), 226–243.
- Love, David A., Michael G. Palumbo, and Paul A. Smith**, “The Trajectory of Wealth in Retirement,” *Journal of Public Economics*, 2009, *93* (1-2), 191–208.
- Marshall, Samuel, Kathleen M. McGarry, and Jonathan S. Skinner**, “The Risk of Out-of-Pocket Health Care Expenditure at End of Life,” 2010. NBER Working Paper 16170.
- Nakajima, Makoto**, “Optimal Capital Income Taxation with Housing,” 2010. Federal Reserve Bank Working Paper No. 10-11.

- **and Irina A. Telyukova**, “Reverse Mortgage Loans: A Quantitative Analysis,” *Journal of Finance*, 2017, *72* (2), 911–950.
- **and** –, “Medical Expenses and Saving in Retirement: The Case of U.S. and Sweden,” 2018. Federal Reserve Bank of Minneapolis Opportunity and Inclusive Growth Institute Working Paper No. 8.
- Ortalo-Magné and Sven Rady**, “Housing Market Dynamics: On the Contribution of Income Shocks and Credit Constraints,” *Review of Economic Studies*, 2006, *73* (2), 459–485.
- Skinner, Jonathan**, “Comment on: Aging and Housing Equity: Another Look by Venti and Wise,” in David A. Wise, ed., *Perspectives on the Economics of Aging*, Chicago: University of Chicago Press, 2004, chapter 3, pp. 176–180.
- Venti, Steven F. and David A. Wise**, “Aging and Housing Equity: Another Look,” in David A. Wise, ed., *Perspectives on the Economics of Aging*, Chicago: University of Chicago Press, 2004, chapter 3, pp. 127–175.
- Yogo, Motohiro**, “Portfolio Choice in Retirement: Health Risk and the Demand for Annuities, Housing, and Risky Assets,” *Journal of Monetary Economics*, 2016, *80*, 17–34.

Appendix (Not for Publication)

Appendix A Data Analysis Robustness: Weighting and Panel Balancing

In this section, we demonstrate how our choices regarding the weight scheme that we use, as well as the way we treat panel balancing, impact the data facts. To remind the reader, in our data analysis, we chose to use 1996 weights for the households in the sample, which means that the households in our analysis have to be present in the first wave, and that thereafter, we are looking at an unbalanced panel. The choice of first-wave weights was motivated by the fact that we do not want to lose nursing-home residents from our sample, while the unbalanced panel is the most natural mapping of the model to the data, since in the model, we will also generate an unbalanced panel, with realistic mortality rates conditioned on all the state variables of the model.




Figure A1: Median Housing Assets: Weight Comparison.




Figure A2: Median Financial Assets: Weight Comparison.

First, we compute the cohort profiles of median housing and financial assets using alternative weighting schemes. Figure A1 and figure A2 compare profiles of median housing and financial assets, with (i) the baseline assumptions (using 1996 weights on the sample, with no subsequent cohort replenishment, labeled wave-3 weights in the graphs), (ii) the same sample, with no sample weighting, and (iii) cohort replenishment and using wave-specific weights, which implies losing nursing home residents. We check the case without sample weighting to compare our results with De Nardi et al. (2010), who do not use sample weighting in their data analysis; our results align with theirs well, given that they use only singles in their analysis, while we also use couples. The pictures that we found under the baseline assumptions, that is, upward-sloping housing asset profiles for all cohorts, which reflect the house price boom during the sample period, and

approximately flat financial asset profiles, are roughly maintained under alternative assumptions. Using weights – either 1996 or cohort-specific – elevates the levels of assets, especially for younger cohorts. We prefer to use the weighted sample, but in such a way that it still allows us to account for nursing home residents with positive weights.




Figure A3: Median Housing Assets: Balanced vs Unbalanced Panel.




Figure A4: Median Financial Assets: Balanced vs Unbalanced Panel.

We also demonstrate the impact of choosing to work with the entire sample (those who were present in 1996), which creates an unbalanced panel, versus working with only those who start in 1996 and survive into the eighth wave of the survey (a balanced panel). Figures A3 and A4 plot median housing and financial asset profiles in the balanced and unbalanced panels. We find that especially for financial assets, using the balanced panel makes the asset profiles steeper, so that asset decumulation over the life cycle looks more pronounced. This confirms what De Nardi et al. (2010) called the mortality bias: including non-survivors in the sample alters the sample composition toward those in poorer health, who also tend to have less wealth, so that the median profiles look flatter as a result.

Appendix B Parameter Sensitivity Experiments

B.1 Role of Housing and Parameter Identification

In this section, we show a series of graphs (figures A5 - A9) for each sensitivity experiment, with the first panel showing the target of the alternative calibration of the parameter, marked with “(T)”, and the other panels showing aspects of the model fit that deteriorate as a result of the perturbation. Then, table A1 gives the quantitative role of housing for each of the experiments.




Figure A5: Sensitivity Experiment – $\underline{c} = 11,600$ to improve renters' Medicaid match




Figure A6: Sensitivity Experiment – $\omega_1 = 1.5$ to improve homeownership rate match

B.2 Bequest and Precautionary Motives

As described in the text, we run two experiments: in both, we raise σ by 25% to 3.8, which is the point estimate in De Nardi et al. (2010), then re-estimate first γ and then β to match only the median net worth profiles. We can reproduce the net worth profile in each period, although as figures A10 and A11 show, the model deteriorates along other target dimensions. In the γ

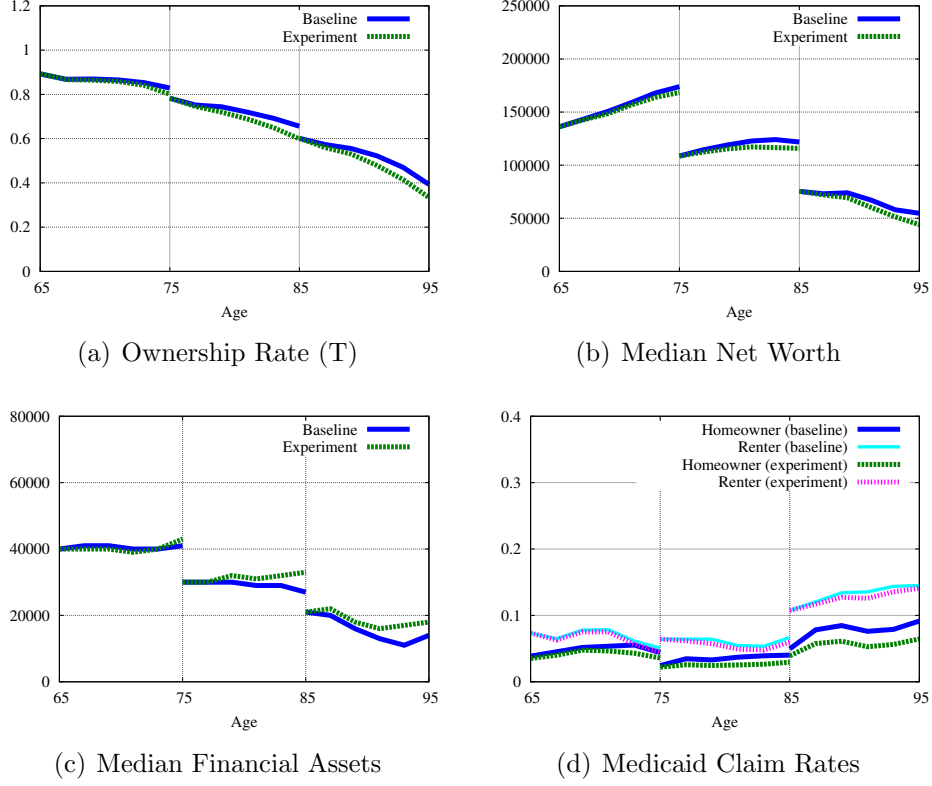




Figure A7: Sensitivity Experiment – $\eta = 0.86$ to improve homeownership rate match

case, this is the rate at which retirees decumulate assets and the median financial asset profile; in the β case, it is the median financial asset and debt rate profiles that suffer. Notice that in both cases, the homeownership rate profile is retained.


Tables A2 and A3 gives the quantitative decomposition of each of the salient model channels. We observe similar interactions to what we saw in the previous experiments. Notably, in spite of the changing relative roles of medical expenses and bequest motives, the role of housing overall remains robust, and the role of individual housing channels (not shown) is equally robust.



(a) Median Financial Assets(T)




(b) Median Net Worth




(c) Owners v. Renters


Figure A8: Sensitivity Experiment – $\beta = 0.99$ to improve median financial asset match




(a) Median Financial Assets(T)



(b) Owners v. Renters




(c) Ownership Rate




(d) Debt Rate


Figure A9: Sensitivity Experiment – $\zeta = 100,000$ to improve median financial asset match




(a) Median Total Assets



(b) Owners v. Renters




(c) Ownership Rate




(d) Median Financial Assets


Figure A10: Sensitivity Experiment – $\sigma = 3.8$, γ reestimated to match median net worth ($\gamma = 0$)




(a) Median Total Assets



(b) Ownership Rate



(c) Debt Rate



(d) Median Financial Assets

Figure A11: Sensitivity Experiment – $\sigma = 3.8$, β reestimated to match median net worth ($\beta = 0.93$)

Table A1: Percent Contribution of Housing to Median Net Worth – Sensitivity Experiments

Parameter change	Age 75 (cohort 1)	Age 85 (cohort 2)	Age 95 (cohort 3)
$\underline{c} = 11,600$	29-41	38-44	53-63
$\omega_1 = 1.5$	26-39	33-38	31-47
$\eta = 0.86$	25-38	30-36	26-42
$\beta = 0.99$	25-36	31-38	33-48
$\zeta = 100,000$	24-38	25-34	22-44

The ranges refer to different orders of introducing housing (before and after bequest motives and medical expenses). See text for the targets that determined alternative parameter values.

Table A2: Contribution of Each Channel to Median Net Worth – Sensitivity Experiment $\sigma = 3.8$, $\gamma = 0$

Model	Age 75 (cohort 1)	Age 85 (cohort 2)	Age 95 (cohort 3)
Simple life-cycle with longevity risk	54.0	49.2	47.0
1 Bequest motive	0.0	0.0	0.0
Total housing	38.1	36.2	31.5
Medical expenditure risk	7.9	14.6	21.5
2 Medical expenditure risk	11.1	13.8	-0.6
Bequest motive	0.0	0.0	0.0
Total housing	34.9	37.0	53.6

Table A3: Contribution of Each Channel to Median Net Worth
– Sensitivity Experiment $\sigma = 3.8, \beta = 0.93$

	Model	Age 75 (cohort 1)	Age 85 (cohort 2)	Age 95 (cohort 3)
	Simple life-cycle with longevity risk	49.4	45.3	42.0
1	Bequest motive	13.0	11.0	8.7
	Total housing	31.6	31.1	27.9
	Medical expenditure risk	6.0	12.6	21.5
2	Total housing	41.1	36.3	31.1
	Bequest motive	3.6	5.9	5.5
	Medical expenditure risk	6.0	12.6	21.5
3	Medical expenditure risk	10.9	15.0	1.7
	Bequest motive	7.4	3.3	7.2
	Total housing	32.3	36.4	49.1
4	Bequest motive	13.0	11.0	8.7
	Medical expenditure risk	5.3	7.3	0.2
	Total housing	32.3	36.4	49.1